

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/305696291>

GPU-based Approaches for Shape Diameter Function Computation and Its Applications Focused on Skeleton Extraction

Article in *Computers & Graphics* · July 2016

DOI: 10.1016/j.cag.2016.06.006

CITATIONS

6

READS

1,225

7 authors, including:



Martin Madaras
Comenius University Bratislava

29 PUBLICATIONS 29 CITATIONS

[SEE PROFILE](#)



Adam Riečický
Comenius University Bratislava

7 PUBLICATIONS 14 CITATIONS

[SEE PROFILE](#)



Roman Ďuríkovič
Comenius University Bratislava

108 PUBLICATIONS 455 CITATIONS

[SEE PROFILE](#)



Andrea Baldacci
Italian National Research Council

5 PUBLICATIONS 9 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



Optical-Inertial Hybrid Motion Capture Sync System [View project](#)



The Neptune Fountain in Bologna (Italy) [View project](#)

GPU-based Approaches for Shape Diameter Function Computation and Its Applications Focused on Skeleton Extraction

Abstract

In this paper two approaches for the computation of the shape diameter function (SDF) on the GPU are outlined and compared. The SDF is a scalar function describing the local thickness of an object. It can be used for consistent mesh partitioning and skeletonization. In the first approach, we have reorganized the tracing of the rays to be well suited for the rasterization hardware. To the best of our knowledge, this is the first method to show how to compute the SDF using only the rasterization hardware and without the need of any acceleration data structures. The second approach uses parallel ray casting and an octree traversal using OpenCL. We demonstrate that the first method achieves similar results as the ray casting using OpenCL. In addition, it is faster for large meshes and it is simpler to implement. Furthermore, we extend the SDF computation by fast post-processing using texture-space diffusion. The fast SDF computation can be used in many applications such as the automatic skeleton extraction as we demonstrate in the article.

Keywords:

shape diameter function, OpenCL, depth peeling, skeleton extraction, skeleton texture mapping

1. Introduction

The shape diameter function (SDF) is a scalar function defined on the mesh surface which expresses a measure of the diameter of the object's volume in the neighborhood of each point on the surface. See Figure 1 for some examples. The original idea was proposed by Shapira et al. [1] and it is related to the concept of Medial Axis Transform (MAT) [2].

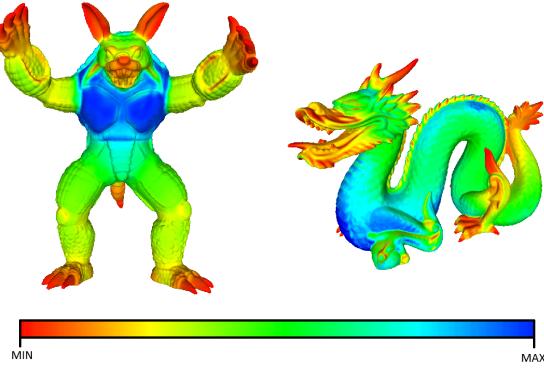


Figure 1: The shape diameter function (SDF) expresses a measure of the local thickness around each point on the surface. The SDF values are colored in a scale from low (red) to high (blue).

The SDF computation is based on creating a cone centered around the inward-normal direction of every point and sending several rays inside this cone to the other side of the mesh (see Figure 3), measuring the distance at the point of intersection. This process, despite being highly parallelizable, is computationally expensive. If processed on the CPU, the task becomes extremely inefficient. Therefore, this paper presents two

algorithms to accelerate the calculation of the SDF. In the first approach, instead of tracing multiple rays per vertex at once, we launch one ray for each vertex multiple times. In fact, our method is based on iteratively peeling two or three successive depth layers of the mesh from multiple views around the mesh. From the review of the current literature, to the best of our knowledge, we can state that our depth peeling approach is the first one that computes the SDF exploiting only the rasterization hardware and without the need of any acceleration data structures. In the second approach, we propose a parallel method for computing the SDF that is performed on the GPU using OpenCL, exploiting the independence of rays. Our approach starts by transferring the data containing indices of triangles and our acceleration structure to the GPU. We then calculate a diameter value for each ray independently. Finally, the computation of the final SDF value is performed for each vertex gathering its computed rays diameters.

Given the above tools for SDF computation, we illustrate a few practical applications of the fast SDF computation. The SDF values have to be computed repeatedly during an iterative process, thus we have to compute the SDF as fast as possible. Proposed applications are the computation of skeleton sheets for mesh parameterization and an improvement of mesh contraction weights to extract finer skeletons with small scale details. Additionally, we illustrate how the SDF can be used for the fast extraction of skeletons with linear topology.

2. Related Work

Shape diameter function The shape diameter function (SDF) was introduced by Shapira et al. [1]. Given a point on the mesh surface, the SDF is defined as the weighted average of

45 the lengths of rays traced from the given point inside the cone
 46 centered at its inward-normal. For each ray, the normal at the
 47 point of intersection is checked against the normal at the start-
 48 ing point. If the angle between the two normals is less than 90
 49 degrees the ray is rejected. From a practical point of view, this
 50 test avoids taking into account rays pointing directly to the out-
 51 side of the mesh. To make this method robust and also immune
 52 to non-watertight meshes, outliers removal is performed pre-
 53 serving only those rays whose lengths fall within one standard
 54 deviation from the median of all lengths. In [3], it is pointed out
 55 that the outliers removal technique proposed in [1] generates
 56 counter-intuitive results in some cases. For example, in the case
 57 of a mesh composed of two parallel (infinite) planes, the SDF
 58 value obtained by the Shapira et al. method is given by the aver-
 59 age of one correct value that is the length of the ray cast along
 60 the normal and the lengths of all the other rays thrown inside
 61 the cone that systematically overestimates the correct diameter.
 62 The same problem occurs at the bifurcation of a Y-shape. In
 63 this case, the correct value would be the minimum of all rays
 64 lengths, thus increasing the opening of the cone has the only
 65 effect of adding noise to the diameter estimation when taking
 66 an average ray length. In [3], in order to resolve the dilemma
 67 between a small or large cone, a more conservative estimation
 68 of the SDF is introduced by using an adaptive cone size. In
 69 particular, the algorithm starts with a large cone aperture that
 70 gets decreased as long as the measure of the absolute growth of
 71 the minimum ray length remains within a certain threshold. For
 72 example, in the case of the infinite parallel planes this method
 73 converges to a small cone size giving a correct SDF value equiv-
 74 alent to the distance between the two planes. Another applica-
 75 tion of SDF was introduced in [4], where SDF is used for mesh
 76 segmentation based on strokes drawn by the user. Essentially,
 77 the main operation in the SDF calculation is given by tracing
 78 rays. In [1], such operation is performed on the CPU using an
 79 octree as an acceleration structure. As Table 1 shows, ray trac-
 80 ing on the CPU is costly and the need for an acceleration struc-
 81 ture is mandatory. A fast approximation of the SDF is proposed
 82 in [5] where the SDF value is computed only for some Poisson-
 83 distributed points over the mesh surface and then the computed
 84 values are propagated across the whole mesh using the Poisson
 85 equation. The latter paper reports only the timing for a single
 86 low poly model of a horse where a performance in the order
 87 of seconds is achieved but at the cost of some approximation.
 88 Instead, our GPU method based on depth peeling computes the
 89 SDF values for each vertex without approximations and with-
 90 out the need of acceleration structures. Moreover, we achieve
 91 performance in the order of few seconds even for meshes with
 92 hundreds of thousands of vertices.

93 **Depth peeling** The depth peeling technique, introduced in
 94 [6], starts by rendering the scene getting the nearest fragment to
 95 the camera. In the subsequent passes, the previous depth buffer
 96 is bound to a texture in a fragment shader. The scene is ren-
 97 dered again and all fragments whose depth is less or equal to
 98 the corresponding depth from the previous pass are discarded.
 99 To avoid read-modify-write hazards, the technique ping-pongs
 100 the depth values between two different buffers. Depth peeling
 101 is a robust image-based solution to different kinds of problems,

102 specifically: order-independent transparency, layered depth im-
 103 age and ray tracing. The most severe drawback of this technique
 104 is that it requires the geometry to be drawn multiple times caus-
 105 ing a strong limitation for GPU-bound applications. A more ef-
 106 ficient implementation of the depth peeling algorithm has been
 107 proposed in [7].

108 **Parallel ray casting** There are several ways to perform ray
 109 casting. Carr et al. [8] proposed to generate rays and perform
 110 traversal of acceleration structures on the CPU, then store the
 111 results and perform the ray-triangle intersections on the GPU.
 112 Purcell el al. [9] proposed to store the scene geometry and the
 113 acceleration structure on the GPU and to use the same device to
 114 perform both traversal and intersection. Recent efforts to opti-
 115 mize the algorithms for GPUs have demonstrated to obtain bet-
 116 ter results on traversal of acceleration structures than a single-
 117 core execution on CPU. Therefore, the second option is more
 118 suitable for the SDF algorithm.

119 3. GPU-based Computation of the Shape Diameter Func- 120 tion

121 We propose two different methods for computing the SDF
 122 using the GPU. The first approach exploits the OpenGL pro-
 123 grammable pipeline via shaders. The second approach relies on
 124 OpenCL to parallelize ray casting for each vertex. We provide
 125 a comparison of two methods in terms of their computed values
 126 and performances.

127 3.1. Parallelized Ray Casting using Depth Peeling

128 The SDF calculation using depth peeling exploits the inher-
 129 ently parallel task of tracing independent rays in a different way
 130 than standard ray tracing. In particular, we start by generating
 131 a set of cameras by uniformly sampling a sphere around the 3D
 132 model and placing an orthogonal camera at each sampled posi-
 133 tion. The camera view direction is set to look down along the
 134 sphere normal at each sampled point. We call the set of cameras
 135 \mathbf{C} . As a preprocessing step, we also pack all vertices in a sin-
 136 gle texture that we call \mathbf{T} . Now, for each camera, we alternate
 137 drawing front or back faces performing a depth peeling until no
 138 more fragments are written or after a maximum amount of iter-
 139 ations is reached. We keep track of the result of the last 3 render
 140 operations using a circular array of 3 framebuffers: \mathbf{F}_0 , \mathbf{F}_1 and
 141 \mathbf{F}_2 . At every odd iteration, we bind the last three generated
 142 depth textures (two at the first pass) and we perform the com-
 143 putation of one diameter value along the current view direction.
 144 We accumulate the diameters computed from all directions in a
 145 texture array. From this array, we calculate the final SDF value
 146 by first removing outliers and then by performing a weighted
 147 sum of the inliers. As in the original algorithm, the weight of
 148 each ray is given by the inverse of the angle between the ray and
 149 the center of the cone around the inward-normal of each vertex.

150 To simplify the comprehension of the algorithm, we refer to
 151 Figure 2 during our explanation. An outline of this method is
 152 reported in Algorithm 1 and Algorithm 2. Given a camera $c \in$
 153 \mathbf{C} , we bind \mathbf{F}_0 and we render the front-most part of the model
 154 as depicted in Figure 2(a). We write to both the color and the

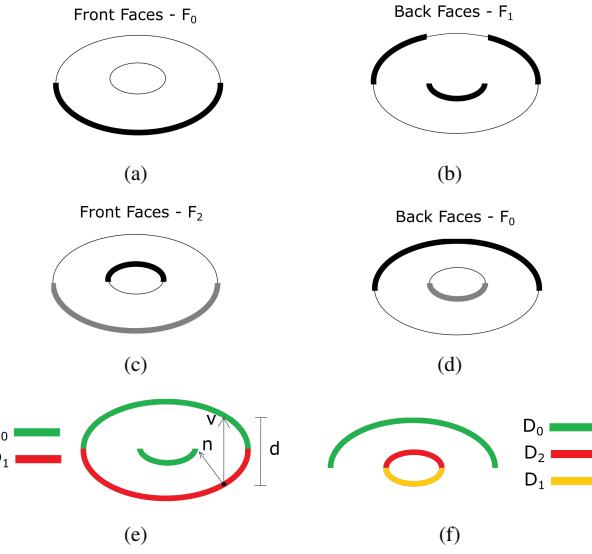


Figure 2: Figures (a), (b), (c) and (d) show the first four iterations of our depth peeling ray tracing applied to a torus. The images show the front-most surfaces as bold black lines, hidden surfaces as thin black lines, and peeled away surfaces as light grey lines. Figures (e) and (f) illustrate the diameter calculation launched at each odd iteration.

depth buffer, using a 32 bit floating point texture for the RGB color channel and a 32 bit depth buffer. In the color channel we store the normals of the vertices. At the next iteration, we draw the back faces of the model to \mathbf{F}_1 as depicted in Figure 2(b). Given \mathbf{F}_0 and \mathbf{F}_1 , we can now calculate the diameter value for the vertices that belong to the first depth layer \mathbf{D}_0 , which is the one stored in \mathbf{F}_0 . The depth layer \mathbf{D}_0 is highlighted in red in Figure 2(e). We bind a framebuffer of the same dimension of the texture \mathbf{T} and we draw a full-screen quad to launch a fragment program for each vertex. The fragment program fetches the position for each vertex from \mathbf{T} . The vertex position is transformed by the current camera matrix $c \in \mathbf{C}$ and its projection in the image space \mathbf{v}_p is found along with its distance \mathbf{v}_z with respect to the camera. At this point, we obtain the depth values \mathbf{z}_0 and \mathbf{z}_1 by sampling respectively the textures \mathbf{D}_0 and \mathbf{D}_1 (red and green lines in Figure 2(e), respectively) at the position \mathbf{v}_p . If $\mathbf{v}_z < \mathbf{z}_1$, we conclude that the vertex must belong to the first depth layer \mathbf{D}_0 and thus its diameter value along this view direction is given by: $d = \text{abs}(\mathbf{v}_z - \mathbf{z}_1)$. Otherwise, if $\mathbf{v}_z > \mathbf{z}_1$ we discard the fragment.

In the third iteration, we render to \mathbf{F}_2 the front faces of the 3D model but discarding all fragments whose depth is less than the one stored in \mathbf{D}_1 . See Figure 2(c) for the resulting depth buffer \mathbf{D}_2 . At the fourth iteration, we render to \mathbf{F}_0 (overwriting all data previously stored) the back faces of the 3D model discarding all fragments whose depth is less than the one stored in \mathbf{D}_2 . See Figure 2(d) for the resulting depth buffer \mathbf{D}_0 . At this point, we can calculate the diameter value for the depth layer \mathbf{D}_2 depicted in red in Figure 2(f). The computation of the diameter value proceeds exactly as in the previous pass except for the depth comparison performed between successive layers. As a matter of fact, at this pass we obtain the depth values \mathbf{z}_0 , \mathbf{z}_2

and \mathbf{z}_1 by sampling respectively the textures \mathbf{D}_0 , \mathbf{D}_2 and \mathbf{D}_1 at the projected vertex position \mathbf{v}_p , where the order of the values is consistent with the use of our circular array.

Algorithm 1 Depth Peeling Ray Tracing

```

Require: set of cameras C
Require: circular array of 3 framebuffers F
Require: diameters texture array A
Require: maximum iterations N
Ensure: SDF value at each vertex
     $i \leftarrow 0$ 
     $j \leftarrow 0$ 
     $k \leftarrow 0$ 
    for each  $c$  in  $C$  do
        while  $i < N$  do
            setCamera( $c$ )
            if  $i == 0$  then
                bindShader(nopeeling)
            else if  $j > 0$  then
                bindShader(peeling)
                bindDepthTexture( $F[j - 1]$ )
            else
                bindShader(peeling)
                bindDepthTexture( $F[2]$ )
            end if
            if  $i \% 2 == 0$  then
                 $F[j] \leftarrow \text{renderModel(GL_FRONT)}$ 
            else
                 $F[j] \leftarrow \text{renderModel(GL_BACK)}$ 
            end if
            if  $i \% 2 != 0$  then
                if  $i > 0$  then
                     $pB \leftarrow (j + 1)\%3$ 
                     $pF \leftarrow (j == 0)?2:(j - 1)$ 
                     $A[k] \leftarrow \text{calcDiameter}(F[pF], F[j], F[pB])$ 
                else
                     $A[k] \leftarrow \text{calcDiameter}(F[j - 1], F[j], \text{NULL})$ 
                end if
                 $k \leftarrow (k + 1)$ 
            end if
             $j \leftarrow (j + 1)\%3$ 
             $i \leftarrow (i + 1)$ 
        end while
    end for
     $SDF \leftarrow \text{weightedSum}(A)$ 

```

Algorithm 2 Calculate Diameter

```

Require: front faces buffer FRONT
Require: back faces buffer BACK
Require: back faces from previous step PREVBACK
Require: texture pack with vertices normals Tn
Require: texture pack with vertices positions Tv
Require: current view direction D
Ensure: diameter value at current vertex
     $N \leftarrow \text{texelFetch}(Tn, fragCoord.xy)$ 
     $angle \leftarrow \text{acos}(N.\text{dot}(D))$ 
    if  $angle > 60$  then
        discard fragment and return
    end if
     $V \leftarrow \text{texelFetch}(Tv, fragCoord.xy)$ 
     $P \leftarrow \text{project}(V)$ 
    if isFalseIntersection() then
        discard fragment and return
    end if
     $zFront \leftarrow \text{texture2D}(FRONT, P.xy)$ 
     $zBack \leftarrow \text{texture2D}(BACK, P.xy)$ 
     $zPrevBack \leftarrow \text{texture2D}(PREVBACK, P.xy)$ 
    if  $zPrevBack < P.z$  and  $P.z < zBack$  then
        return  $zBack - zFront$ 
    else
        discard fragment and return
    end if

```

190 In the fragment shader for SDF computation, we thus accept
191 only vertices whose depth \mathbf{v}_z satisfies the following inequality:
192 $\mathbf{z}_1 < \mathbf{v}_z < \mathbf{z}_0$. The comparison using the three layers is necessary

because vertices whose depth satisfies the condition $v_z < z_0$ but not the condition $z_1 < v_z$ already got their values computed from the previous passes. The algorithm described proceeds until no more fragments get rasterized or after a fixed number of iterations. To test if any fragment is drawn to the framebuffer, we simply start an occlusion query before each depth peeling iteration. We now point out some implementation details neglected until now for the sake of clarity. When evaluating the SDF function, the corresponding fragment program is early discarded if the current view direction does not form an angle of less than 60 degrees with the inward-normal of the vertex under consideration. Moreover, to perform the false intersection test from [1], we also pack all the vertex normals in a texture and we test the vertex normal against the normal at the point of intersection stored in the color texture drawn during the depth peeling pass. Finally, once a single diameter value is computed for a vertex, it is stored using the *imageStore* command in a texture array bound to the SDF fragment program as an *image2DArray*. Similarly, a counter contained in an unsigned integer texture is increased using the *imageAtomicAdd* function. At the end of the whole process, we have an unsigned integer texture containing the number of rays traced per vertex and a texture array containing all the ray lengths for each vertex. To gather the final SDF value, we bind a framebuffer of the same dimension of the texture T and we draw a fullscreen quad. For each vertex, we now loop over the diameter values contained in the texture array and we sort them using a naive bubble sort. Then, we calculate mean and variance of the values and we calculate the weighted sum of all rays lengths that are less than one standard deviation away from the median (which is taken as the middle element of the sorted values). We can perform SDF calculation both for vertices and faces. The aforementioned explanation assumes that values are calculated per vertex. For the face case, we trace rays through the barycenter of the triangle.

3.2. Parallelized Ray Casting using OpenCL

Our second approach for the computation of the SDF extends the original implementation of the algorithm by porting it to the GPU using OpenCL. The computation of the SDF is a three-step process that consists of ray generation, ray casting and traversal of the chosen acceleration structure and finally the computation of the distance at the point of intersection.

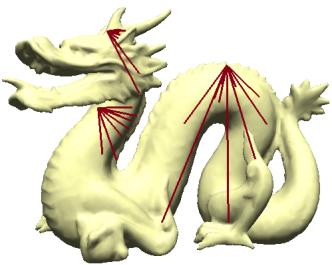


Figure 3: Rays are cast in cones from each vertex. The cast rays are displayed as red lines.

In the first step, we have to generate rays in a cone centered around the inward-normal of a given point and their associated

weights (Figure 3). As in the original algorithm, the weight of each ray is given by the inverse of the angle between the ray and the center of the cone. The generation of rays on the CPU and their transfer to the GPU create unnecessary overhead because we must store every ray and its associated weight. Therefore, the rays should be generated on the GPU. However, random generation of rays would imply having a pseudo-random generator in the OpenCL. Rolland [3] has tackled this problem by defining a cone sampling strategy consisting of random rays that are uniformly generated inside the cone. Our algorithm further expands the former method by evenly distributing the rays in a cone using the spherical Fibonacci point set [10]. Furthermore, a uniform distribution of rays helps prevent unnecessary bias in the values on the mesh. A comparison of random ray generation with respect to our approach can be seen in Figure 4. The algorithm generates rays restricted to a given sphere cap.

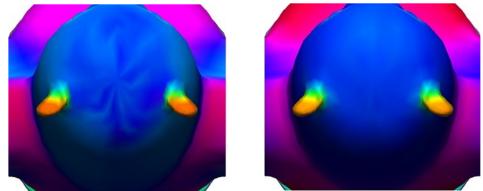


Figure 4: (Left) random ray generation, (right) uniform ray generation using spherical Fibonacci point set.

Each ray is then transformed into the world coordinates by multiplying it by the tangent space matrix specified by the point's inward-normal and by the two orthogonal unit vectors spanning the tangent plane of the point. It is important to note that only valid points are used to generate the rays, where a point is valid if it has a non-zero normal, tangent and binormal vectors. In the case of non-manifold models, the use of triangle barycenters instead of vertices is more reliable as the normal of some points cannot be properly determined. As a side effect, this also helps reduce unnecessary data transfer thanks to the fact that normal, tangent and binormal can be easily calculated.

In the second step, for each ray, we traverse our acceleration structure. The acceleration structure is one of the most important parts of ray tracing. For the purpose of comparison with the original article, we built an octree structure, even though a BVH or kDTree could be traversed faster. Our octree implementation is based on Laine's stack-based approach [11].

In the third step, we compute ray-triangle intersections between the cast ray and the triangles contained in the found octree node. For this purpose, we use the Moller and Trumbore algorithm [12] because it mainly uses the dot and cross products which can be evaluated very efficiently on current graphics hardware. Once the values per ray are calculated, they are transferred back to the CPU, where the final outliers removal takes place. The rays with lengths that do not fall within one standard deviation from the median of all lengths are removed.

The inliers are then averaged using the weights calculated in the first step. To overcome the errors in the measure caused by pose changes, a smoothing operation is necessary. Therefore, we perform a *k*-ring neighborhood Gaussian smoothing, where

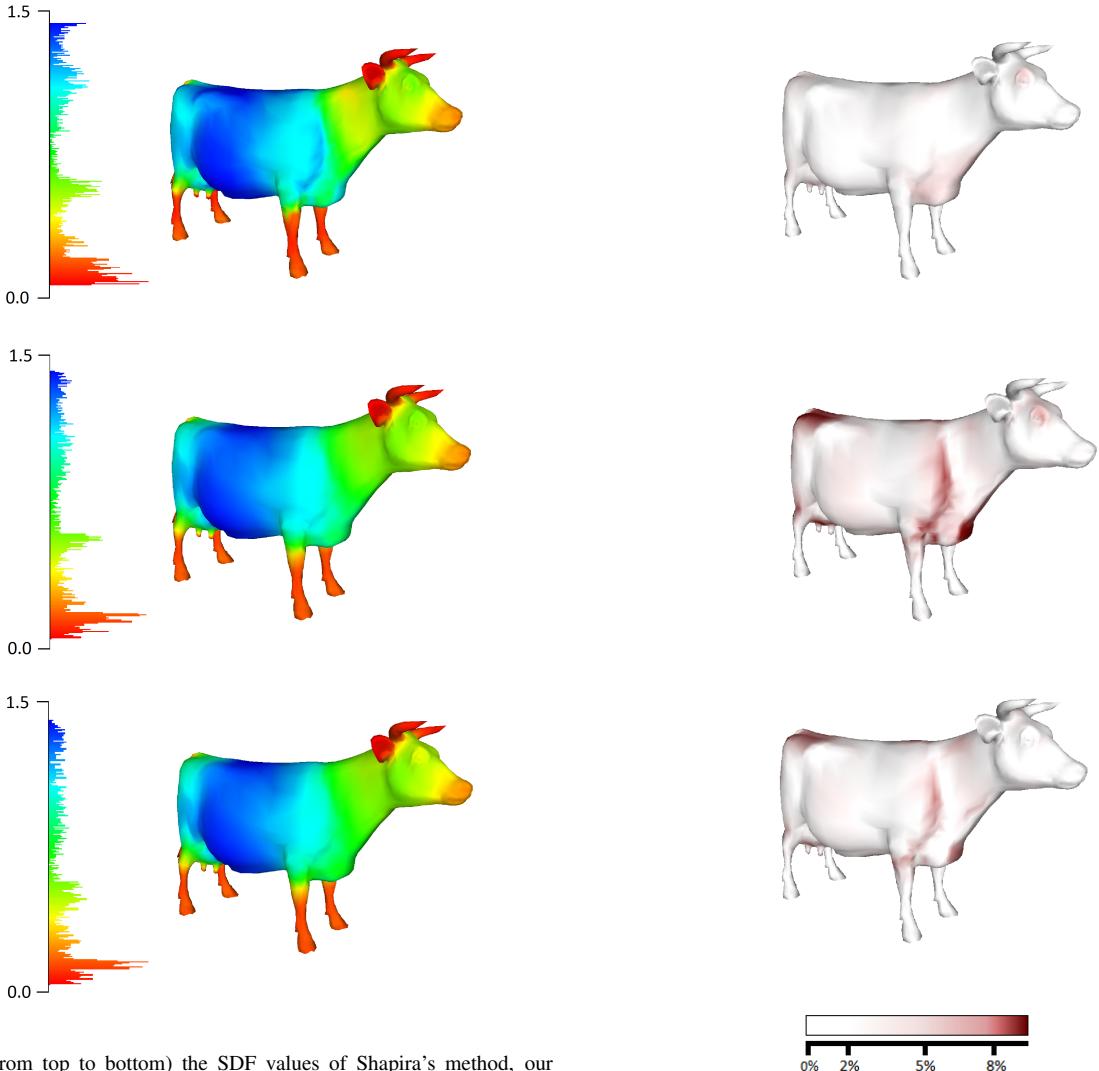


Figure 5: (From top to bottom) the SDF values of Shapira’s method, our OpenGL depth peeling and our OpenCL raycasting. The SDF values are colored in a scale from low (red) to high (blue) and the histogram of computed SDF values is on the left side of the image.

²⁸² k specifies the blur radius in terms of edges in neighbourhood
²⁸³ radius.

²⁸⁴ 4. Results and Comparison

²⁸⁵ We now compare the results from both our methods for par-
²⁸⁶ allel computation of the SDF. Finally, at the end of this section,
²⁸⁷ a texture-space diffusion method is proposed in order to smooth
²⁸⁸ the values across the surface. A performance comparison of the
²⁸⁹ two approaches is presented in Table 1. Furthermore, we have
²⁹⁰ compared the values of OpenGL depth peeling, OpenCL ray-
²⁹¹ casting and the original Shapira’s implementation. The results
²⁹² of all three sets of SDF values with their related histograms
²⁹³ can be seen in Figure 5. The differences among the methods
²⁹⁴ mapped into a red scale are visualized in Figure 6. To ensure
²⁹⁵ the correctness of our implementation, several basic tests have
²⁹⁶ been performed on different shapes with known diameters. Fig-
²⁹⁷ ure 7 shows all three methods applied on a circle swept along
²⁹⁸ a bezier curve with a diameter linearly interpolated between 2

Figure 6: (From top to bottom) the difference of SDF values of OpenGL depth peeling vs. OpenCL raycasting, Shapira vs. OpenCL raycasting and Shapira vs. OpenGL depth peeling. (Bottom) the color scale of differences from white to red.

²⁹⁹ and 20 units. All three methods give almost the same results
³⁰⁰ and the highest deviation (up to 10%) from the real diameter has
³⁰¹ been measured in the thickest parts of the model using the orig-
³⁰² inal Shapira’s method. Tests on all 3D models were performed
³⁰³ using 30 rays for each point and with a cone aperture of 120.
³⁰⁴ For the depth peeling method, we used 256 camera positions
³⁰⁵ around each model. The differences between the SDF values of
³⁰⁶ our two methods can be seen in Figure 9. All tests have been run
³⁰⁷ on a desktop PC with the following configuration: Intel Core i5,
³⁰⁸ 2,67GHz with 4GB RAM and AMD Radeon R9 290. As you
³⁰⁹ can see in Table 1, the OpenCL raycasting is more efficient for
³¹⁰ 3D models with a modest number of vertices. For 3D models
³¹¹ with a high vertex count (above 150k), the depth peeling ap-
³¹² proach is instead definitely faster. The processing times for the
³¹³ CPU approach using Shapira’s method [1] and the processing
³¹⁴ times for the smoothing phase are also presented in Table 1.
³¹⁵ To perform the smoothing, each vertex and his neighbors are

321 **4.1. Postprocessing using Texture-space Diffusion**

322 We can use the extracted skeleton and a skeleton texture
 323 mapping (STM) [13] to blur the SDF values on the surface of
 324 the object. The idea is to map the SDF values into texture space,
 325 make the Gaussian diffusion there and map the smoothed val-
 326 ues back to the model surface, as shown in Figure 10. Texture-

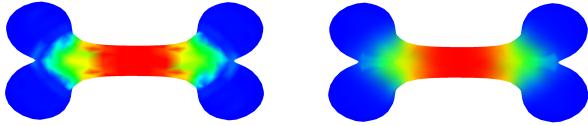


Figure 10: Computed SDF values (left) are stored in the texture and blurred (right) using texture-space diffusion with STM parameterization. Here, a kernel with radius 10px was used.

327 Space diffusion using STM guarantees that local neighborhood
 328 of each point on the model surface is homotopic to local neigh-
 329 borhood in the texture space. Comparison of smoothing on the
 330 mesh using various k -ring areas and smoothing in the texture
 331 using Gaussian filter with various radius settings can be seen in
 332 Table 2.

Radius	k -ring	CPU Smoothing	STM Smoothing
64	6	14,882	0.140
32	5	10,764	0.093
16	4	7,566	0.057
8	3	5,024	0.051
4	2	3,120	0.048
2	1	1,762	0.041

Table 2: Smoothing in a 2048 × 2048 texture using various radius settings.
 Times shown are in seconds [s].

333 **5. Further Applications and Extensions**

334 In this section, we present some applications of computed
 335 SDF values focusing on skeleton extraction. These applications
 336 need to compute SDF values at each step during an iterative
 337 process. Therefore, these applications of SDF are a motivation
 338 for fast parallel SDF computation on the GPU. In the skele-
 339 ton texture mapping approach presented in [13], we need to
 340 compute an internal skeleton from which every mesh vertex is
 341 visible. There are automatic techniques for curve-skeleton ex-
 342 traction [14, 15, 16], but extracted curve-skeletons using these
 343 techniques do not satisfy the above mentioned condition. Be-
 344 cause of this, we have extended the Laplacian-based smoothing
 345 method presented in [14] with an SDF term that moves skeleton
 346 vertices near the medial axis transform (MAT). In addition, the
 347 SDF values can be used to modify parameters as weights and
 348 grouping distances during Laplacian-based skeleton extraction
 349 to obtain better results with skeleton extraction. For further de-
 350 tails, we refer the reader to the original Laplacian-based mesh
 351 contraction [14]. Finally, we show an application of the SDF
 352 for the fast linear skeleton extraction.

353 **5.1. Proposed SDF Modifications for Laplacian-based Skele-
 354 ton Extraction**

355 In order to parameterize a mesh using a skeleton texture
 356 map (STM) [13], we have to guarantee that the skeleton is reli-
 357 able. Reliability refers to the property of the curve-skeleton that
 358 every boundary point (i.e., the point on object’s surface) is vis-
 359 ible from at least one skeleton location [17]. Some basic ideas
 360 behind the modifications of skeleton extraction algorithm are
 361 shown here. However, the in-depth discussion of this topic is
 362 out of the scope of this paper. For a more detailed description,
 363 we refer the reader to [18].

364 **5.1.1. Skeleton Sheets using SDF**

In most cases, a mesh can be parameterized by a curve-
 skeleton, but not in general. As a matter of fact, in some cases,
 the skeleton must contain sheets, similarly to the medial axis
 transform (MAT). A skeleton consisting of a mixture of curves
 and surface sheets was defined in [19] and called a “meso-
 skeleton”. In [19], authors used Voronoi poles to guide the sur-
 face toward the medial axis. Here we show that this is the place
 where the SDF can be efficiently used. Shifted vertices into
 the central sheet using SDF are less noisy than Voronoi poles
 (or MAT) and SDF can be computed faster using our GPU ap-
 proaches. Furthermore, our method creates skeleton sheets only
 in the regions where the curve-skeleton reliability is not satis-
 fied. Meso-skeletons created using Voronoi poles [19] present
 sheet surfaces also in the regions which could be parameterized
 using just the curve-skeletons. This is clearly not efficient for
 STM because skeleton sheet takes much more space in the final
 texture. Therefore, we want to have the skeleton sheets only in
 the regions where they are inevitable. We extended the mesh



Figure 11: Examples of skeleton sheets of C-shape and S-shape models. Classical Laplacian-based skeleton extraction fails here because of non-convex cross-
 sections along the skeleton.

contraction method for skeleton extraction [14] with an SDF centering term as follows:

$$\begin{bmatrix} \mathbf{W}_L \mathbf{L} \\ \mathbf{W}_H \\ \mathbf{W}_C \end{bmatrix} \mathbf{V}' = \begin{bmatrix} \mathbf{0} \\ \mathbf{W}_H \mathbf{V} \\ \mathbf{W}_C \mathbf{C} \end{bmatrix}, \quad (1)$$

365 where matrix \mathbf{C} is composed of vertices shifted by the SDF
 366 values in the opposite of normal direction $\mathbf{c}_i = \mathbf{v}_i - (SDF_i/2) \cdot$
 367 \mathbf{n}_i . \mathbf{v}_i , \mathbf{n}_i and SDF_i are respectively the position, normal and
 368 SDF value of the vertex i . For a more detailed description of

369 the original contraction weights (\mathbf{W}_L , \mathbf{W}_H) and the introduced
 370 weights \mathbf{W}_C we refer reader to [14] and to [18], respectively.

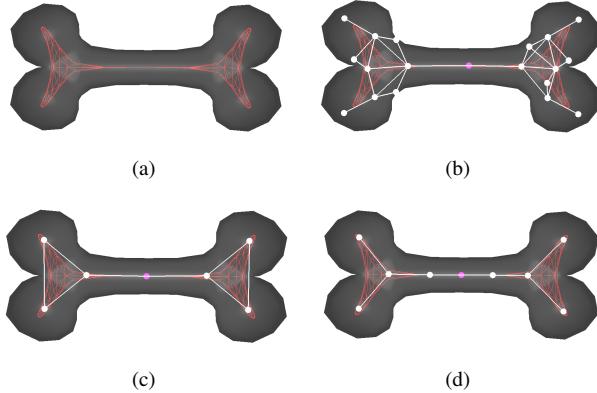
371 During the iterative contraction process, we tend to collapse
 372 the vertices into the medial axis composed of vertices shifted by
 373 half of the SDF values in the opposite of the normal direction.
 374 Using this extended version, we are able to create the curve-
 375 skeleton in the parts of the model where this is reliable and
 376 skeleton sheets where the curve-skeleton would not satisfy the
 377 reliability condition (see Figure 11).

378 5.2. Modification of Iterative Parameters

379 At each iteration of the iterative contraction process, the
 380 SDF values are computed and used to drive the modification
 381 of the contraction weights and grouping distance parameters.
 382 This results in skeletons which reflect better small-scale details
 383 of the 3D model (see Figure 14).

384 5.2.1. Updating of Grouping Distance using SDF

385 Here we show how SDF values can be used to create a skele-
 386 ton from a not fully contracted mesh. During the simplification
 387 of the contracted mesh, we add another condition upon which a
 388 half-edge collapse is applied. This new condition is satisfied if
 389 an Euclidean distance between two nodes is less than a distance
 390 threshold dt_i . The distance threshold for each node is computed
 391 as: $dt_i = ggt \cdot diag + (SDF_i/2) \cdot gmt$, where ggt is a global group-
 392 ing tolerance term, SDF_i is an SDF value for vertex i , gmt is
 393 global multiplication term and $diag$ is a diagonal of the model’s
 394 bounding box. We obtained the best results with values set to
 395 $ggt = 0.05$ and $gmt = 1.25$. These parameters enable all the
 396 uncontracted parts of the mesh to collapse while the contracted
 397 parts of the mesh remain unchanged (see Figure 12).

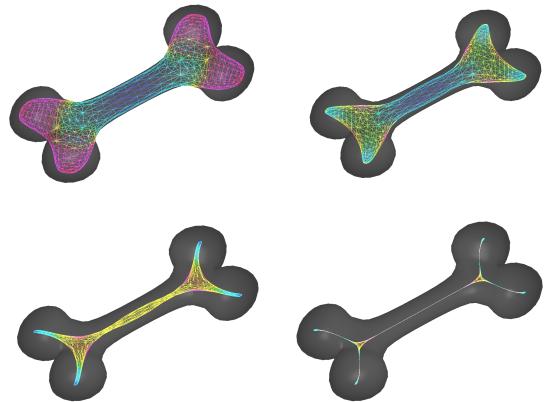


398 Figure 12: Grouping distance using SDF. (a) Contraction of the mesh when
 399 the volume is too big for skeleton extraction, (b) using lower global grouping
 400 distance, (c) using higher global grouping distance, (d) using local grouping
 401 distance based on SDF.

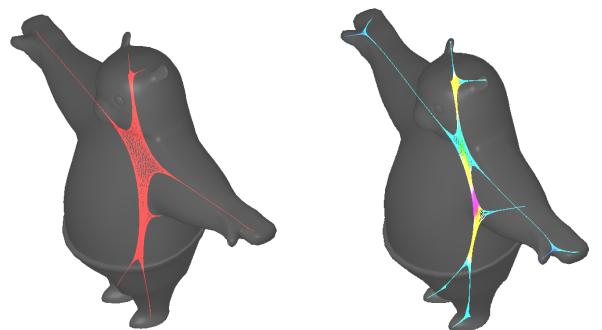
398 5.2.2. Updating of Contraction Weights

400 During the iterative contraction process, we change the hold-
 401 ing weight \mathbf{W}_{Hi} after each iteration based on the SDF values.
 402 We scale \mathbf{W}_{Hi} according to the normalized SDF value $SDFnorm_i$.
 403 First, we update \mathbf{W}_{Hi} according to the surface change of one

403 ring areas as $\mathbf{W}_{Hi}^{t+1} = \mathbf{W}_{Hi}^0 \sqrt{A_i^0/A_i^t}$, where A_i^0 and A_i^t are the
 404 original and current areas of adjacent faces for vertex i , respec-
 405 tively. Then we normalize the SDF values to the interval $\langle 0, 1 \rangle$
 406 as $SDFnorm_i = (SDF_i - minSDF)/(maxSDF - minSDF)$,
 407 where $minSDF$ and $maxSDF$ are the minimal and the maxi-
 408 mal SDF values. Afterwards, we change \mathbf{W}_{Hi} according to the
 409 volume change using SDF as $\mathbf{W}_{Hi} = \mathbf{W}_{Hi}/(SDFnorm_i * (1.0 -$
 410 $bias) + bias)$, where $bias$ is an offset value because we want to
 411 avoid division by zero. We set bias as 0.01 in our tests, which
 412 produces high \mathbf{W}_{Hi} values for $SDFnorm_i = 0$, while avoiding a
 413 division by zero. Using our iterative modification of contraction
 414 weights (Figure 13), we were able to preserve small-scale de-
 415 tails also in cases where the thickness of the 3D model changes
 416 rapidly (see Figure 14). More details concerning this applica-
 417 tion can be found in [18].



418 Figure 13: Contraction weights using SDF. The mesh is colored according to
 419 the SDF values. Red regions with higher SDF values are more contracted and
 420 blue regions with lower SDF value are less contracted.



418 Figure 14: Difference of small-scale details as hands, ears and tail without SDF
 419 (left) and with the SDF term (right).

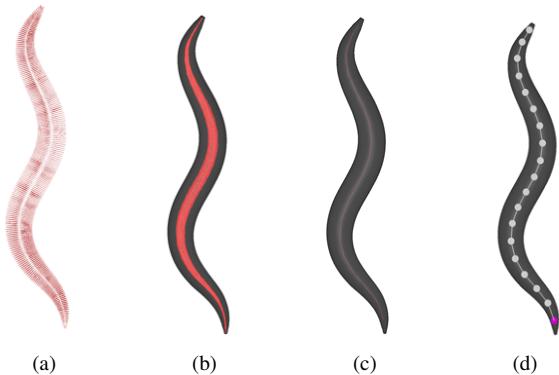
418 5.3. Fast Linear Skeleton Extraction

419 The SDF values have been used for the fast linear skeleton
 420 extraction from particle simulation. We used our fast parallel
 421 GPU implementation for linear skeleton extraction from parti-
 422 cle simulation of the worm body [20]. The extracted skeleton

423 was used to compress the simulation and to visualize it using a
 424 skinning algorithm.

425 In order to compute the SDF of an input point cloud, a surface
 426 must be reconstructed. The reconstructed surface must be
 427 closed, but it must not be 2D-manifold. Therefore, we con-
 428 struct a local Delaunay triangulations for the neighborhood of
 429 each vertex in its tangent plane. In the next step, all edges in
 430 the local triangulations are copied into the global triangulation.
 431 This results in a graph that is a non-manifold orientable closed
 432 global triangulation.

433 Having unoriented point clouds, a normal can be estimated
 434 using PCA. However, the orientation of the normal remains un-
 435 known posing a serious problem for a correct SDF calculation.
 436 Therefore, we orient the normals in two steps. First, we choose
 437 a vertex whose normal orientation can be easily decided. For
 438 this purpose, we choose the outermost vertex in one direction,
 439 for example the topmost vertex. For this vertex, we are able to
 440 choose the normal orientation, because we know that its y coor-
 441 dinate has to be positive, facing out of the mesh bounding box.
 442 We then propagate the orientation over the mesh surface, by
 443 comparing normal orientations between neighboring vertices.
 444 At the end of this process we obtain an oriented point cloud.



485 Figure 15: The algorithm for skeleton extraction from the worm cuticle based
 486 on SDF. (a) Mesh of cuticle, (b) SDF half-vectors. (c) Shifted mesh vertices,
 487 (d) poly line obtained by agglomerative clustering and (e) extracted skeleton
 488 structure.

489 The cuticle of the worm has a simple topology without branch-
 490 ing and a fast skeleton extraction algorithm is needed. The al-
 491 gorithm first computes the SDF values and then constructs the
 492 final skeleton using agglomerative clustering with the distance
 493 threshold t (see Figure 15). From clustered polyline a more
 494 subdivided skeleton can be easily constructed. The method is
 495 faster than all known robust methods and works well for the
 496 worm topology.

497 In order to match skeletons from all the timesteps and cal-
 498 culate the rotations between different timesteps, we need to en-
 499 sure that each skeleton has the same number of segments. Thus,
 500 we find a clustering threshold t using binary search between
 501 the minimal and maximal SDF values computed on the surface.
 502 The computed threshold t is then used for agglomerative clus-
 503 tering during the skeleton construction.

460 6. Discussion

461 The way in which the parallelized OpenCL raycasting pro-
 462 ceeds is similar to Shapira’s approach [1]. The OpenGL depth
 463 peeling method proposes another point of view to approach the
 464 SDF calculation. Both proposed approaches achieved similar
 465 results in average, even though the depth peeling method is a
 466 bit faster in comparison to the parallelized version of ray casting
 467 technique, especially when the input model has more than 150k
 468 vertices. This is due to the fact that the whole SDF processing
 469 is calculated in a buffer space. Therefore, after the rendering
 470 phase, the calculation is independent of the mesh complexity.
 471 Both GPU methods are faster than former method as can be
 472 seen in Table 1. Unfortunately, we were not able to compare
 473 times with the Rolland’s method [3], because computational
 474 times were not mentioned in their paper. The comparison of
 475 our two implementations against the Shapira’s version shows
 476 that the results obtained are very similar. For this test, we set up
 477 our implementations to follow the details outlined in the orig-
 478 inal paper. The differences in the calculated values are very
 479 likely due to implementation details which were not mentioned
 480 in the original paper. Finally, as demonstrated in Figure 8, our
 481 depth peeling method is robust to pose changes with essentially
 482 no differences to the original Shapira’s method with given a suf-
 483 ficient sampling of the view sphere around the model.

484 7. Conclusions and Future Work

485 We have implemented two different approaches for the com-
 486 putation of the SDF and compared their results and processing
 487 times. Both approaches are implemented on the GPU. The first
 488 one is calculated using shaders and depth peeling while the sec-
 489 ond one is based on parallelized ray casting using OpenCL. As
 490 mentioned in the discussion, we found out that it is more effi-
 491 cient to use OpenCL raycasting for relatively small meshes and
 492 OpenGL depth peeling for relatively large ones.

493 Furthermore, we presented several applications of the SDF.
 494 We have extended Laplacian smoothing-based skeleton extrac-
 495 tion by an SDF term to obtain skeletons with finer details. We
 496 also extended curve skeletons to skeleton sheets which are con-
 497 tained in models with non-convex cross-sections. Finally, we
 498 used our customized SDF linear skeleton extraction for the fast
 499 compression of particle simulation based on skinning.

500 In both of our GPU implementations for the SDF compu-
 501 tation, we maintain the original outliers removal approach pro-
 502 posed in [1], leaving the one proposed by [3] for future work.
 503 As a matter of fact, the adaptive cone size proposed in [3] helps
 504 with the estimation of the SDF only in some special cases, but
 505 it is computationally more expensive than original method as
 506 the ray casting must be iterated several times to find the correct
 507 cone size.

508 References

- 509 [1] L. Shapira, A. Shamir, D. Cohen-Or, Consistent mesh partitioning and
 510 skeletonisation using the shape diameter function, Vis. Comput. 24 (2008)
 511 249–259.

- 512 [2] H. I. Choi, S. W. Choi, H. P. Moon, Moon: Mathematical theory of medial
 513 axis transform, *Pacific J. Math.*
- 514 [3] X. Rolland-Nevière, G. Doërr, P. Alliez, Robust diameter-based thickness
 515 estimation of 3d objects, *Graphical Models* 75 (6) (2013) 279–296.
- 516 [4] L. Fan, L. Lic, K. Liu, Paint mesh cutting, *Computer Graphics Forum*
 517 30 (2) (2011) 603–612.
- 518 [5] M. Kovacic, F. Guggeri, S. Marras, R. Scateni, Fast approximation of the
 519 shape diameter function, in: *GraVisMa 2010 Proceedings*, 2010, p. 1724.
- 520 [6] C. Everitt, Interactive order-independent transparency, White paper,
 521 nVIDIA 2 (6) (2001) 7.
- 522 [7] L. Bavoil, K. Myers, Order independent transparency with dual depth
 523 peeling, Tech. rep., NVIDIA Developer SDK 10 (Feb. 2008).
- 524 [8] N. A. Carr, J. D. Hall, J. C. Hart, The ray engine, in: *Proceedings of*
 525 the ACM SIGGRAPH/EUROGRAPHICS Conference on Graphics Hard-
 526 ware, HWWS '02, Eurographics Association, Aire-la-Ville, Switzerland,
 527 Switzerland, 2002, pp. 37–46.
- 528 [9] T. J. Purcell, I. Buck, W. R. Mark, P. Hanrahan, Ray tracing on pro-
 529 grammable graphics hardware, *ACM Trans. Graph.* 21 (3) (2002) 703–
 530 712.
- 531 [10] R. Marques, C. Bouville, M. Ribardire, L. P. Santos, K. Bouatouch,
 532 Spherical fibonacci point sets for illumination integrals, *Computer Graph-
 533 ics Forum* 32 (8) (2013) 134–143.
- 534 [11] S. Laine, T. Karras, Efficient sparse voxel octrees, in: *Proceedings of*
 535 the 2010 ACM SIGGRAPH Symposium on Interactive 3D Graphics and
 536 Games, I3D '10, ACM, New York, NY, USA, 2010, pp. 55–63.
- 537 [12] T. Möller, B. Trumbore, Fast, minimum storage ray-triangle intersection,
 538 *J. Graph. Tools* 2 (1) (1997) 21–28.
- 539 [13] M. Madaras, R. Őurikovič, Skeleton texture mapping, in: *Proceedings*
 540 of the 28th Spring Conference on Computer Graphics, SCCG '12, ACM,
 541 New York, NY, USA, 2013, pp. 121–127.
- 542 [14] O. K.-C. Au, C.-L. Tai, H.-K. Chu, D. Cohen-Or, T.-Y. Lee, Skeleton
 543 extraction by mesh contraction, *ACM Trans. Graph.* 27 (3) (2008) 44:1–
 544 44:10.
- 545 [15] G. Aujay, F. Hétroy, F. Lazarus, C. Depraz, Harmonic skeleton for re-
 546 alistic character animation, in: *Proceedings of the 2007 ACM SIG-
 547 GRAPH/Eurographics Symposium on Computer Animation, SCA '07*,
 548 Eurographics Association, Aire-la-Ville, Switzerland, Switzerland, 2007,
 549 pp. 151–160.
- 550 [16] T. K. Dey, J. Sun, Defining and computing curve-skeletons with medial
 551 geodesic function, in: *Proceedings of the 4th EG symposium on Geom.
 552 processing*, 2006, pp. 143–152.
- 553 [17] N. D. Cornea, D. Silver, P. Min, Curve-skeleton properties, applica-
 554 tions, and algorithms, *IEEE Transactions on Visualization and Computer
 555 Graphics* 13 (3) (2007) 530–548.
- 556 [18] R. Őurikovič, M. Madaras, Mathematical Progress in Expressive Im-
 557 age Synthesis II: Extended and Selected Results from the Symposium
 558 MEIS2014, Springer Japan, Tokyo, 2015, Ch. Controllable Skeleton-
 559 Sheets Representation Via Shape Diameter Function, pp. 79–90.
- 560 [19] A. Tagliasacchi, I. Alhashim, M. Olson, H. Zhang, Mean curvature skele-
 561 tons, *Comp. Graph. Forum* 31 (5) (2012) 1735–1744.
- 562 [20] M. Piovarči, M. Madaras, R. Őurikovič, Physically inspired stretching
 563 for skinning animation of non-rigid bodies, in: *Proceedings of the 31st
 564 Spring Conference on Computer Graphics, SCCG '15*, ACM, New York,
 565 NY, USA, 2015, pp. 47–53.